

# StockMarketPredictionUsingML

November 5, 2022

## 1 Visualizing and Forecasting of Stocks (NSE)

- Pandas – This library helps to load the data frame in a 2D array format and has multiple functions to perform analysis tasks in one go.
- Numpy – Numpy arrays are very fast and can perform large computations in a very short time.
- Matplotlib/Seaborn – This library is used to draw visualizations.
- Sklearn – This module contains multiple libraries having pre-implemented functions to perform tasks from data preprocessing to model development and evaluation.
- XGBoost – This contains the eXtreme Gradient Boosting machine learning algorithm which is one of the algorithms which helps us to achieve high accuracy on predictions.

```
[1]: #importing the libraries
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import pandas as pd
from pandas.plotting import lag_plot
import glob
import os
sns.set()
import warnings
warnings.filterwarnings('ignore')

import chart_studio.plotly as py
import plotly.graph_objs as go
from plotly.offline import plot
import seaborn as sb
#for offline plotting
from plotly.offline import download_plotlyjs, init_notebook_mode, plot, iplot
init_notebook_mode(connected=True)

import seaborn as sb
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from xgboost import XGBClassifier
from sklearn import metrics
```

## 1.1 Importing Dataset

- The dataset we will use here to perform the analysis and build a predictive model is Adani Enterprise Stock Price data.
- We will use OHLC('Open', 'High', 'Low', 'Close') data from 1st January 2019 to 3rd Nov 2022

```
[2]: #loading the data
adanient = pd.read_csv('NSE_ADANIENT.csv', parse_dates=['Date'], dayfirst=True)
#adanient['Date'] = pd.to_datetime(adanient['Date'])
adanient.head()
```

```
[2]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2019-01-01	160.899994	162.350006	155.449997	157.250000	155.465088	
1	2019-01-02	157.000000	157.850006	152.500000	154.850006	153.092361	
2	2019-01-03	154.899994	156.100006	150.300003	152.500000	150.769012	
3	2019-01-04	152.100006	154.000000	150.000000	152.550003	150.818451	
4	2019-01-07	152.899994	154.449997	150.250000	151.250000	149.533203	

	Volume
0	4726656
1	2735262
2	2758876
3	2777308
4	2714218

- From the first five rows, we can see that data for some of the dates is missing the reason for that is on weekends and holidays Stock Market remains closed hence no trading happens on these days.

```
[3]: adanient.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 950 entries, 0 to 949
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        950 non-null   datetime64[ns]
1   Open        950 non-null   float64
2   High        950 non-null   float64
3   Low         950 non-null   float64
4   Close       950 non-null   float64
5   Adj Close   950 non-null   float64
6   Volume      950 non-null   int64
dtypes: datetime64[ns](1), float64(5), int64(1)
memory usage: 52.1 KB
```

```
[4]: print("\n")
```

```

print("Open    --- mean :", np.mean(adanient['Open']), " \t Std: ", np.
↳std(adanient['Open']), " \t Max: ", np.max(adanient['Open']), " \t Min:␣
↳", np.min(adanient['Open']))
print("High    --- mean :", np.mean(adanient['High']), " \t Std: ", np.
↳std(adanient['High']), " \t Max: ", np.max(adanient['High']), " \t Min:␣
↳", np.min(adanient['High']))
print("Low     --- mean :", np.mean(adanient['Low']), " \t Std: ", np.
↳std(adanient['Low']), " \t Max: ", np.max(adanient['Low']), " \t Min:␣
↳", np.min(adanient['Low']))
print("Close   --- mean :", np.mean(adanient['Close']), " \t Std: ", np.
↳std(adanient['Close']), " \t Max: ", np.max(adanient['Close']), " \t Min:␣
↳", np.min(adanient['Close']))
print("Volume  --- mean :", np.mean(adanient['Volume']), " \t Std: ", np.
↳std(adanient['Volume']), " \t Max: ", np.max(adanient['Volume']), " \t Min:␣
↳", np.min(adanient['Volume']))

```

Open	---	mean :	967.6365255178948	Std:	959.4566600451187	Max:
3837.649902		Min:	116.349998			
High	---	mean :	986.2366821631572	Std:	975.2937875729492	Max:
3885.0		Min:	119.5			
Low	---	mean :	949.9067905957897	Std:	943.5604917335505	Max:
3812.0		Min:	113.0			
Close	---	mean :	969.0835262031584	Std:	960.392715442575	Max:
3834.550049		Min:	116.949997			
Volume	---	mean :	4898751.874736842	Std:	5184593.824054264	Max:
61334483		Min:	248249			

```
[5]: adanient.shape
```

```
[5]: (950, 7)
```

- From this, we got to know that there are 950 rows of data available and for each row, we have 7 different features or columns.

```
[6]: adanient['Date'] = pd.to_datetime(adanient['Date'])
```

```

[7]: print(f'Dataframe contains stock prices between {adanient.Date.min()} {adanient.
↳Date.max()}')
print(f'Total days = {(adanient.Date.max() - adanient.Date.min()).days} days')

```

Dataframe contains stock prices between 2019-01-01 00:00:00 2022-11-03 00:00:00  
Total days = 1402 days

```
[8]: adanient.describe()
```

```
[8]:
```

	Open	High	Low	Close	Adj Close \
count	950.000000	950.000000	950.000000	950.000000	950.000000
mean	967.636526	986.236682	949.906791	969.083526	968.184123
std	959.962036	975.807506	944.057495	960.898585	961.141005
min	116.349998	119.500000	113.000000	116.949997	115.622528
25%	154.000000	157.037506	150.250000	153.162495	152.074284
50%	437.100006	454.399994	418.725006	446.600006	446.114136
75%	1632.562500	1669.462463	1608.412537	1644.287476	1643.588806
max	3837.649902	3885.000000	3812.000000	3834.550049	3834.550049

	Volume
count	9.500000e+02
mean	4.898752e+06
std	5.187325e+06
min	2.482490e+05
25%	2.156603e+06
50%	3.549359e+06
75%	5.559366e+06
max	6.133448e+07

```
[9]: adanient.isnull().sum()
```

```
[9]: Date          0
      Open          0
      High          0
      Low           0
      Close         0
      Adj Close     0
      Volume        0
      dtype: int64
```

```
[10]: adanient = adanient.dropna()
```

```
[11]: adanient.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 950 entries, 0 to 949
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        950 non-null   datetime64[ns]
1   Open        950 non-null   float64
2   High        950 non-null   float64
3   Low         950 non-null   float64
4   Close       950 non-null   float64
5   Adj Close   950 non-null   float64
6   Volume      950 non-null   int64
dtypes: datetime64[ns](1), float64(5), int64(1)
```

memory usage: 52.1 KB

```
[12]: adanient.shape
```

```
[12]: (950, 7)
```

- checking the the null values if any are present in the data frame.

```
[13]: adanient.isnull().sum()
```

```
[13]: Date          0
      Open          0
      High          0
      Low           0
      Close         0
      Adj Close     0
      Volume        0
      dtype: int64
```

- This implies that there are no null values in the data set provided.

## 1.2 Exploratory Data Analysis

- EDA is an approach to analyzing the data using visual techniques. It is used to discover trends, and patterns, or to check assumptions with the help of statistical summaries and graphical representations.
- While performing the EDA of the Adani Enterprise Stock Price data we will analyze how prices of the stock have moved over the period of time and how the end of the quarters affects the prices of the stock.


```
[14]: plt.figure(figsize=(20,5))
      plt.plot(adanient['Close'])
      plt.title('Adani Ent Close price.', fontsize=15)
      plt.ylabel('Price in Rupees.')
      plt.show()
```



The prices of the Adani Ent stocks are showing an upward trend as depicted by the plot of the closing price of the stocks


```
[15]: adanient[['Open', 'High', 'Low', 'Close', 'Adj Close']].plot(kind='box')
```

```
[15]: <AxesSubplot:>
```




```
[16]: adanient['High'].plot(figsize=(16,6))
```

```
[16]: <AxesSubplot:>
```



```
[17]: adanient=adanient.set_index('Date',drop=True) # yarra added we can remove later_
      ↪for testing purpose
      ## xlimit and y limit
      adanient['High'].plot(xlim=['2019-01-01','2022-11-04'],figsize=(12,4))
```

```
[17]: <AxesSubplot:xlabel='Date'>
```



```
[18]: adanient.index
```

```
[18]: DatetimeIndex(['2019-01-01', '2019-01-02', '2019-01-03', '2019-01-04',
                    '2019-01-07', '2019-01-08', '2019-01-09', '2019-01-10',
                    '2019-01-11', '2019-01-14',
                    ...,
                    '2022-10-20', '2022-10-21', '2022-10-24', '2022-10-25',
                    '2022-10-27', '2022-10-28', '2022-10-31', '2022-11-01',
                    '2022-11-02', '2022-11-03'],
                    dtype='datetime64[ns]', name='Date', length=950, freq=None)
```

```
[19]: index=adanient.loc['2019-01-01':'2022-11-01'].index
      share_open=adanient.loc['2019-01-01':'2022-11-01']['Open']
```

```
[20]: share_open
```

```
[20]: Date
2019-01-01    160.899994
2019-01-02    157.000000
2019-01-03    154.899994
2019-01-04    152.100006
2019-01-07    152.899994
...
2022-10-25    3321.949951
2022-10-27    3319.850098
2022-10-28    3322.000000
```


```
2022-10-31    3339.949951
2022-11-01    3361.899902
Name: Open, Length: 948, dtype: float64
```

```
[21]: index
```

```
[21]: DatetimeIndex(['2019-01-01', '2019-01-02', '2019-01-03', '2019-01-04',
                  '2019-01-07', '2019-01-08', '2019-01-09', '2019-01-10',
                  '2019-01-11', '2019-01-14',
                  ...,
                  '2022-10-18', '2022-10-19', '2022-10-20', '2022-10-21',
                  '2022-10-24', '2022-10-25', '2022-10-27', '2022-10-28',
                  '2022-10-31', '2022-11-01'],
                  dtype='datetime64[ns]', name='Date', length=948, freq=None)
```

```
[22]: figure,axis=plt.subplots()
      plt.tight_layout()
      ## Preventing overlapping
      figure.autofmt_xdate()
      axis.plot(index,share_open)
```

```
[22]: [<matplotlib.lines.Line2D at 0x18e11c7c6a0>]
```






```
[23]: adanient[['Close']].plot()  
plt.title("Adani Ent")  
plt.show()
```




```
[24]: # Cumulative Return  
dr = adanient.cumsum()  
dr.plot()  
plt.title('Adani Ent Cumulative Returns')
```

```
[24]: Text(0.5, 1.0, 'Adani Ent Cumulative Returns')
```



```
[25]: plt.figure(figsize=(10,10))
lag_plot(adanient['Open'], lag=5)
plt.title('Adani Ent Autocorrelation plot')
plt.show()
```

\*c\* argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with \*x\* & \*y\*. Please use the \*color\* keyword-argument or provide a 2D array with a single row if you intend to specify the same RGB or RGBA value for all points.



```
[26]: adanient=adanient.reset_index()
```

```
[27]: adanient.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 950 entries, 0 to 949
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        950 non-null   datetime64[ns]
1   Open        950 non-null   float64
2   High        950 non-null   float64
3   Low         950 non-null   float64
4   Close       950 non-null   float64
5   Adj Close   950 non-null   float64
```

```

6    Volume      950 non-null    int64
dtypes: datetime64[ns](1), float64(5), int64(1)
memory usage: 52.1 KB

```

```
[28]: adanient.set_index('Date',drop=True)
```

```
[29]: adanient.head()
```

```
[29]:
```

	Open	High	Low	Close	Adj Close \
Date					
2019-01-01	160.899994	162.350006	155.449997	157.250000	155.465088
2019-01-02	157.000000	157.850006	152.500000	154.850006	153.092361
2019-01-03	154.899994	156.100006	150.300003	152.500000	150.769012
2019-01-04	152.100006	154.000000	150.000000	152.550003	150.818451
2019-01-07	152.899994	154.449997	150.250000	151.250000	149.533203

```

Volume
Date
2019-01-01    4726656
2019-01-02    2735262
2019-01-03    2758876
2019-01-04    2777308
2019-01-07    2714218

```

```
[30]: ## datetime
from datetime import datetime
```

### 1.3 Time Resampling

```
[31]: adanient.resample(rule='A').min()
```

```
[31]:
```

	Open	High	Low	Close	Adj Close \
Date					
2019-12-31	116.349998	119.500000	113.000000	116.949997	115.622528
2020-12-31	121.000000	129.800003	116.400002	120.900002	120.768463
2021-12-31	477.000000	493.250000	477.000000	490.899994	490.365906
2022-12-31	1574.900024	1592.000000	1528.800049	1543.949951	1543.293945

```

Volume
Date
2019-12-31    1003411
2020-12-31     620753
2021-12-31     272261
2022-12-31     248249

```

```
[32]: adanient.resample(rule='A').max()
```


```
[32]:
```

	Open	High	Low	Close	Adj Close \
Date					
2019-12-31	217.750000	221.500000	212.399994	218.800003	216.974716
2020-12-31	492.000000	507.000000	484.700012	490.850006	490.315948
2021-12-31	1762.949951	1788.900024	1735.550049	1763.050049	1762.300903
2022-12-31	3837.649902	3885.000000	3812.000000	3834.550049	3834.550049

	Volume
Date	
2019-12-31	61334483
2020-12-31	49264537
2021-12-31	43530006
2022-12-31	15060223


```
[33]: ##year end frequency
adanient.resample(rule='A').max()['Open'].plot()
plt.show()
```



- We are downsample the data using the alias “A” for year-end frequency for Open column and plot results in plot

```
[34]: ##quarterly start frequency
##https://towardsdatascience.com/resample-function-of-pandas-79b17ec82a78
adanient.resample(rule='QS').max()['High'].plot()
```

[34]: <AxesSubplot:xlabel='Date'>



- We are downsample the data using the alias “A” for Quaterly starting data frequency for High column and plot results in plot

```
[35]: ##Business End Frequency  
##https://towardsdatascience.com/resample-function-of-pandas-79b17ec82a78  
adanient.resample(rule='BA').max()
```

```
[35]:
```

	Open	High	Low	Close	Adj Close	\
Date						
2019-12-31	217.750000	221.500000	212.399994	218.800003	216.974716	
2020-12-31	492.000000	507.000000	484.700012	490.850006	490.315948	
2021-12-31	1762.949951	1788.900024	1735.550049	1763.050049	1762.300903	
2022-12-30	3837.649902	3885.000000	3812.000000	3834.550049	3834.550049	

	Volume
Date	
2019-12-31	61334483
2020-12-31	49264537
2021-12-31	43530006
2022-12-30	15060223

```
[36]: adanient.resample(rule='BQS').max()
```

```
[36]:
```


	Open	High	Low	Close	Adj Close \
Date					
2019-01-01	160.899994	162.350006	155.449997	157.250000	155.465088
2019-04-01	169.500000	180.800003	157.699997	161.050003	159.221970
2019-07-01	160.100006	161.800003	154.300003	156.100006	154.328156
2019-10-01	217.750000	221.500000	212.399994	218.800003	216.974716
2020-01-01	256.899994	261.000000	252.649994	258.649994	256.492279
2020-04-01	164.000000	169.199997	161.300003	162.149994	161.973587
2020-07-01	301.799988	322.399994	295.100006	307.549988	307.215393
2020-10-01	492.000000	507.000000	484.700012	490.850006	490.315948
2021-01-01	1063.000000	1093.000000	1018.400024	1058.400024	1057.248535
2021-04-01	1715.000000	1717.199951	1613.000000	1700.949951	1699.099365
2021-07-01	1594.900024	1628.449951	1566.050049	1587.599976	1586.925415
2021-10-01	1762.949951	1788.900024	1735.550049	1763.050049	1762.300903
2022-01-03	2005.000000	2042.000000	1991.000000	2014.750000	2013.893921
2022-04-01	2404.949951	2420.949951	2336.300049	2395.300049	2394.282227
2022-07-01	3837.649902	3885.000000	3812.000000	3834.550049	3834.550049
2022-10-03	3580.000000	3625.149902	3542.449951	3590.399902	3590.399902

	Volume
Date	
2019-01-01	14695085
2019-04-01	61334483
2019-07-01	15687573
2019-10-01	14162661
2020-01-01	10810744
2020-04-01	16533100
2020-07-01	49264537
2020-10-01	20009492
2021-01-01	31098225
2021-04-01	43530006
2021-07-01	14250256
2021-10-01	7093354
2022-01-03	5126019
2022-04-01	6515102
2022-07-01	15060223
2022-10-03	7578847

```
[37]: ##plotting
adanient['Open'].resample(rule='A').mean().plot(kind='bar')
```


```
[37]: <AxesSubplot:xlabel='Date'>
```



```
[38]: adanient['Open'].resample(rule='M').max().plot(kind='bar',figsize=(15,6))
```

```
[38]: <AxesSubplot:xlabel='Date'>
```





```
[39]: adanient['High'].rolling(11).max().head(20)
```

```
[39]: Date
2019-01-01      NaN
2019-01-02      NaN
2019-01-03      NaN
2019-01-04      NaN
2019-01-07      NaN
2019-01-08      NaN
2019-01-09      NaN
2019-01-10      NaN
2019-01-11      NaN
2019-01-14      NaN
2019-01-15    162.350006
2019-01-16    157.850006
2019-01-17    156.100006
2019-01-18    154.449997
2019-01-21    154.449997
2019-01-22    154.449997
2019-01-23    154.449997
2019-01-24    152.600006
2019-01-25    152.600006
2019-01-28    151.250000
Name: High, dtype: float64
```

```
[40]: adanient.head()
```

```
[40]:
```

	Open	High	Low	Close	Adj Close	\
Date						
2019-01-01	160.899994	162.350006	155.449997	157.250000	155.465088	
2019-01-02	157.000000	157.850006	152.500000	154.850006	153.092361	
2019-01-03	154.899994	156.100006	150.300003	152.500000	150.769012	
2019-01-04	152.100006	154.000000	150.000000	152.550003	150.818451	
2019-01-07	152.899994	154.449997	150.250000	151.250000	149.533203	

```

Volume
Date
2019-01-01  4726656
2019-01-02  2735262
2019-01-03  2758876
2019-01-04  2777308
2019-01-07  2714218

```

```
[41]: adanient['Open:30 days rolling']=adanient['Open'].rolling(30).mean()
```

```
[42]: adanient.head(31)
```

```
[42]:
```


	Open	High	Low	Close	Adj Close	\
Date						
2019-01-01	160.899994	162.350006	155.449997	157.250000	155.465088	
2019-01-02	157.000000	157.850006	152.500000	154.850006	153.092361	
2019-01-03	154.899994	156.100006	150.300003	152.500000	150.769012	
2019-01-04	152.100006	154.000000	150.000000	152.550003	150.818451	
2019-01-07	152.899994	154.449997	150.250000	151.250000	149.533203	
2019-01-08	150.500000	154.300003	149.649994	153.449997	151.708237	
2019-01-09	153.600006	154.449997	147.600006	149.350006	147.654785	
2019-01-10	151.000000	151.800003	148.500000	151.100006	149.384918	
2019-01-11	151.100006	152.600006	149.100006	151.500000	149.780380	
2019-01-14	151.000000	151.149994	147.000000	148.350006	146.666122	
2019-01-15	148.500000	149.949997	147.250000	148.649994	146.962708	
2019-01-16	148.199997	150.699997	148.000000	148.449997	146.764984	
2019-01-17	148.399994	151.250000	146.350006	149.000000	147.308746	
2019-01-18	149.699997	149.800003	144.550003	146.750000	145.084274	
2019-01-21	146.500000	148.000000	144.250000	145.949997	144.293365	
2019-01-22	145.399994	146.649994	141.699997	145.699997	144.046219	
2019-01-23	145.500000	146.500000	142.050003	143.000000	141.376846	
2019-01-24	143.000000	144.050003	141.100006	142.550003	140.931961	
2019-01-25	142.850006	145.449997	136.149994	137.350006	135.790985	
2019-01-28	138.000000	141.000000	118.949997	136.399994	134.851761	
2019-01-29	136.449997	143.800003	133.600006	140.800003	139.201828	
2019-01-30	140.899994	141.500000	136.500000	139.899994	138.312027	
2019-01-31	139.500000	141.750000	135.000000	137.149994	135.593246	
2019-02-01	136.949997	144.800003	134.550003	143.000000	141.376846	
2019-02-04	142.899994	142.899994	121.300003	123.099998	121.702728	

2019-02-05	122.000000	127.400002	119.599998	124.250000	122.839676
2019-02-06	123.849998	124.199997	114.800003	122.900002	121.505005
2019-02-07	122.199997	126.800003	119.050003	123.949997	122.543083
2019-02-08	123.000000	125.449997	115.699997	123.750000	122.345345
2019-02-11	122.949997	123.250000	115.800003	118.199997	116.858345
2019-02-12	117.199997	120.000000	115.349998	116.949997	115.622528

	Volume	Open:30 days rolling
Date		
2019-01-01	4726656	NaN
2019-01-02	2735262	NaN
2019-01-03	2758876	NaN
2019-01-04	2777308	NaN
2019-01-07	2714218	NaN
2019-01-08	2791866	NaN
2019-01-09	4054809	NaN
2019-01-10	3534546	NaN
2019-01-11	2254999	NaN
2019-01-14	2624582	NaN
2019-01-15	1805678	NaN
2019-01-16	1574176	NaN
2019-01-17	2618859	NaN
2019-01-18	2306520	NaN
2019-01-21	1763205	NaN
2019-01-22	2261667	NaN
2019-01-23	2322990	NaN
2019-01-24	1678843	NaN
2019-01-25	3297651	NaN
2019-01-28	9404485	NaN
2019-01-29	4301493	NaN
2019-01-30	3102926	NaN
2019-01-31	3498642	NaN
2019-02-01	7091209	NaN
2019-02-04	9832462	NaN
2019-02-05	8484607	NaN
2019-02-06	9905219	NaN
2019-02-07	9826247	NaN
2019-02-08	8059813	NaN
2019-02-11	7048335	143.391665
2019-02-12	3698931	141.934999

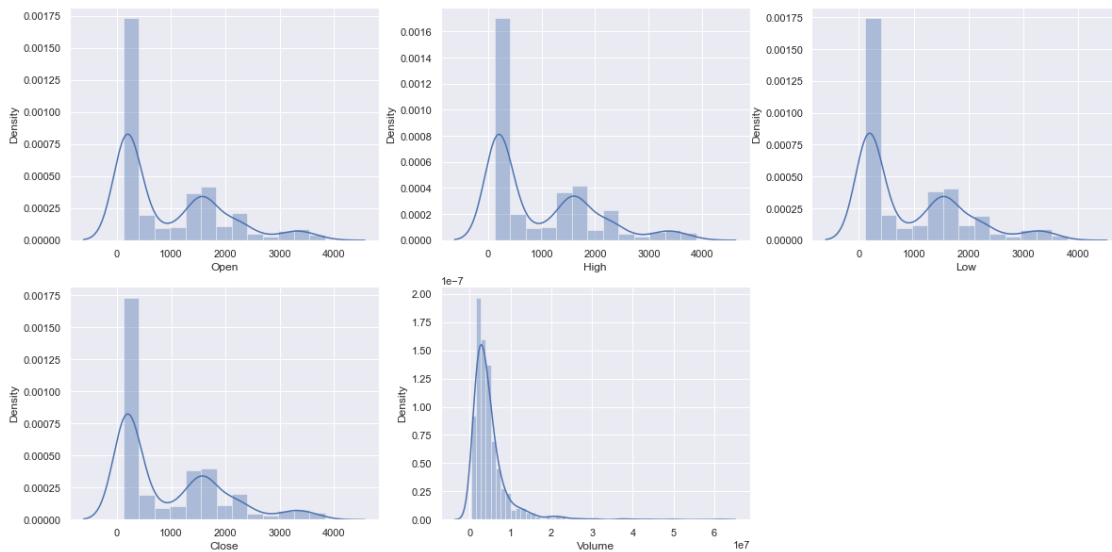
```
[43]: adanient[['Open', 'Open:30 days rolling']].plot(figsize=(12,5))
```

```
[43]: <AxesSubplot:xlabel='Date'>
```



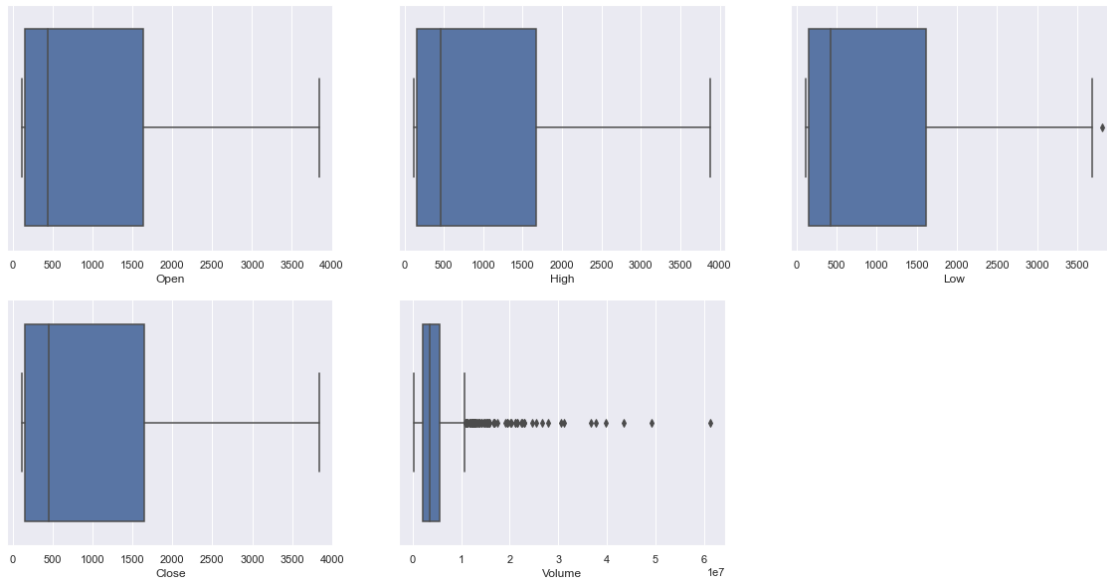
[ ]:

```
[44]: features = ['Open', 'High', 'Low', 'Close', 'Volume']
plt.subplots(figsize=(20,10))
for i, col in enumerate(features):
    plt.subplot(2,3,i+1)
    sb.distplot(adanient[col])
plt.show()
```



- In the distribution plot of OHLC data, we can see two peaks which means the data has varied significantly in two regions. And the Volume data is left-skewed.

```
[45]: plt.subplots(figsize=(20,10))
      for i, col in enumerate(features):
          plt.subplot(2,3,i+1)
          sb.boxplot(adanient[col])
      plt.show()
```



- From the above boxplots, we can conclude that only volume data contains outliers in it but the data in the rest of the columns are free from any outlier.

## 1.4 Feature Engineering

```
[46]: adanient=adanient.reset_index()
```

```
[47]: adanient['year'] = adanient['Date'].dt.year
      adanient['month'] = adanient['Date'].dt.month
      adanient['day'] = adanient['Date'].dt.day
      adanient.head()
```

```
[47]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2019-01-01	160.899994	162.350006	155.449997	157.250000	155.465088	
1	2019-01-02	157.000000	157.850006	152.500000	154.850006	153.092361	
2	2019-01-03	154.899994	156.100006	150.300003	152.500000	150.769012	
3	2019-01-04	152.100006	154.000000	150.000000	152.550003	150.818451	
4	2019-01-07	152.899994	154.449997	150.250000	151.250000	149.533203	

	Volume	Open:30 days rolling	year	month	day
0	4726656	NaN	2019	1	1
1	2735262	NaN	2019	1	2

2	2758876	NaN	2019	1	3
3	2777308	NaN	2019	1	4
4	2714218	NaN	2019	1	7

```
[48]: # set index using column
adanient = adanient.set_index('Date')
```

```
[49]: adanient['is_quarter_end'] = np.where(adanient['month']%3==0,1,0)
adanient.head()
```

```
[49]:
```

	Open	High	Low	Close	Adj Close \
Date					
2019-01-01	160.899994	162.350006	155.449997	157.250000	155.465088
2019-01-02	157.000000	157.850006	152.500000	154.850006	153.092361
2019-01-03	154.899994	156.100006	150.300003	152.500000	150.769012
2019-01-04	152.100006	154.000000	150.000000	152.550003	150.818451
2019-01-07	152.899994	154.449997	150.250000	151.250000	149.533203

	Volume	Open:30 days rolling	year	month	day	is_quarter_end	
Date							
2019-01-01	4726656		NaN	2019	1	1	0
2019-01-02	2735262		NaN	2019	1	2	0
2019-01-03	2758876		NaN	2019	1	3	0
2019-01-04	2777308		NaN	2019	1	4	0
2019-01-07	2714218		NaN	2019	1	7	0

```
[50]: adanientPivot = pd.pivot_table(adanient, values = "Close", columns = "year",
↳ index = "month")
```


```
[51]: adanientPivot.head()
```

```
[51]:
```


year	2019	2020	2021	2022
month				
1	147.121739	218.793480	517.032498	1776.404993
2	124.339474	242.610524	707.327505	1713.949988
3	140.208334	161.490476	933.876194	1771.292864
4	141.044736	139.463888	1155.536840	2217.150024
5	137.509090	138.813157	1293.480005	2156.814261

```
[52]: adanientPivot.plot()
```

```
[52]: <AxesSubplot:xlabel='month'>
```




```
[53]: adanientPivot.plot(subplots = True, figsize=(15, 15), layout=(4,4), sharey=True)
plt.show()
```



```
[54]: data_grouped = adanient.groupby('year').mean()
plt.subplots(figsize=(20,20))

for i, col in enumerate(['Open', 'High', 'Low', 'Close']):
    plt.subplot(2,2,i+1)
    plt.ylabel('Price', fontsize=15)
    plt.title(col, fontsize=20)
    data_grouped[col].plot.bar()
plt.show()
```



- From the above bar graph, we can conclude that the stock prices have 15 times from the year 2019 to that in 2022.

```
[55]: adanient.groupby('is_quarter_end').mean()
```

```
[55]:
```

	Open	High	Low	Close \
is_quarter_end				
0	941.272633	959.315850	924.765064	943.065141
1	1020.531170	1040.248731	1000.349369	1021.284970

	Adj Close	Volume	Open:30 days rolling	year \
is_quarter_end				
0	942.172049	4.778843e+06	935.303652	2020.443218



1	1020.372904	5.139329e+06	957.697521	2020.443038
---	-------------	--------------	------------	-------------

	month	day
is_quarter_end		
0	5.862776	15.733438
1	7.281646	15.762658


Here are some of the important observations of the above-grouped data:

- Prices are higher in the months which are quarter end as compared to that of the non-quarter end months.
- The volume of trades is lower in the months which are quarter end.

```
[56]: adanient['open-close'] = adanient['Open'] - adanient['Close']
      adanient['low-high'] = adanient['Low'] - adanient['High']
      adanient['target'] = np.where(adanient['Close'].shift(-1) > adanient['Close'], 1, 0)
```


- Above we have added some more columns which will help in the training of our model. We have added the target feature which is a signal whether to buy or not we will train our model to predict this only. But before proceeding let's check whether the target is balanced or not using a pie chart.

```
[57]: plt.figure(figsize=(10, 10))
      plt.pie(adanient['target'].value_counts().values,
              labels=[0, 1], autopct='%1.1f%%')
      plt.show()
```



- When we add features to our dataset we have to ensure that there are no highly correlated features as they do not help in the learning process of the algorithm.

```
[58]: plt.figure(figsize=(10, 10))  
  
# As our concern is with the highly  
# correlated features only so, we will visualize  
# our heatmap as per that criteria only.  
sb.heatmap(adanient.corr() > 0.9, annot=True, cbar=False)  
plt.show()
```



- From the above heatmap, we can say that there is a high correlation between OHLC that is pretty obvious and the added features are not highly correlated with each other or previously provided features which means that we are good to go and build our model.

## 1.5 Data Splitting and Normalization

```
[59]: features = adanient[['open-close', 'low-high', 'is_quarter_end']]
target = adanient['target']

scaler = StandardScaler()
features = scaler.fit_transform(features)

X_train, X_valid, Y_train, Y_valid = train_test_split(
    features, target, test_size=0.1, random_state=2022)
```

```
print(X_train.shape, X_valid.shape)
```

```
(855, 3) (95, 3)
```

- After selecting the features to train the model on we should normalize the data because normalized data leads to stable and fast training of the model. After that whole data has been split into two parts with a 90/10 ratio so, that we can evaluate the performance of our model on unseen data.

## 1.6 Model Development and Evaluation

- Now is the time to train some state-of-the-art machine learning models(Logistic Regression, Support Vector Machine, XGBClassifier), and then based on their performance on the training and validation data we will choose which ML model is serving the purpose at hand better.

```
[60]: models = [LogisticRegression(), SVC(
        kernel='poly', probability=True), XGBClassifier()]

for i in range(3):
    models[i].fit(X_train, Y_train)

    print(f'{models[i]} : ')
    print('Training Accuracy : ', metrics.roc_auc_score(Y_train, models[i].
    ↪predict_proba(X_train)[:,-1]))
    print('Validation Accuracy : ', metrics.roc_auc_score(Y_valid, models[i].
    ↪predict_proba(X_valid)[:,-1]))
    print()
```

```
LogisticRegression() :
Training Accuracy :  0.5689750692520776
Validation Accuracy :  0.4316712834718375
```


```
SVC(kernel='poly', probability=True) :
Training Accuracy :  0.5809085872576177
Validation Accuracy :  0.44182825484764543
```

```
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
               colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
               early_stopping_rounds=None, enable_categorical=False,
               eval_metric=None, feature_types=None, gamma=0, gpu_id=-1,
               grow_policy='depthwise', importance_type=None,
               interaction_constraints='', learning_rate=0.300000012,
               max_bin=256, max_cat_threshold=64, max_cat_to_onehot=4,
               max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
               missing=nan, monotone_constraints='()', n_estimators=100,
               n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,
               ...) :
Training Accuracy :  0.9893351800554017
Validation Accuracy :  0.4764542936288088
```

- Among the three models, we have trained XGBClassifier has the highest performance but it is prone to overfitting as the difference between the training and the validation accuracy is too high. But in the case of the Logistic Regression, this is not the case.
- Now let's plot a confusion matrix for the validation data.

```
[61]: plt.figure(figsize=(20, 20))
      metrics.plot_confusion_matrix(models[0], X_valid, Y_valid)
      plt.show()
```

<Figure size 1440x1440 with 0 Axes>



## 1.7 Model Linear Regression model

```
[62]: adanient=adanient.reset_index()
      # Setting the layout for our plot
      layout = go.Layout(
          title='Stock Prices of Adani Enterprises',
          xaxis=dict(
              title='Date',
              titlefont=dict(
                  family='Courier New, monospace',
                  size=18,
                  color='#7f7f7f'
```

```

    )
),
yaxis=dict(
    title='Price',
    titlefont=dict(
        family='Courier New, monospace',
        size=18,
        color='#7f7f7f'
    )
)
)

adanient_data = [{'x':adanient['Date'], 'y':adanient['Close']}]
plot = go.Figure(data=adanient_data, layout=layout)

```

```
[63]: #plot(plot) #plotting offline
      iplot(plot)
```

```
[64]: # Building the regression model
      from sklearn.model_selection import train_test_split

      #For preprocessing
      from sklearn.preprocessing import MinMaxScaler
      from sklearn.preprocessing import StandardScaler
      from sklearn.linear_model import LinearRegression

      import chart_studio.plotly as py
      import plotly.graph_objs as go
      from plotly.offline import plot

      #for offline plotting
      from plotly.offline import download_plotlyjs, init_notebook_mode, plot, iplot
      init_notebook_mode(connected=True)

      #For model evaluation
      from sklearn.metrics import mean_squared_error as mse
      from sklearn.metrics import r2_score

```

```
[65]: #Split the data into train and test sets
      X = np.array(adanient.index).reshape(-1,1)
      Y = adanient['Close']
      X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3,
      ↪random_state=101)
```

```
[66]: # Feature scaling
      scaler = StandardScaler().fit(X_train)
```

```
[ ]:
```

```
[67]: #Creating a linear model
lm = LinearRegression()
lm.fit(X_train, Y_train)
```

```
[67]: LinearRegression()
```

```
[68]: #Plot actual and predicted values for train dataset
trace0 = go.Scatter(x = X_train.T[0],y = Y_train,mode = 'markers',name = 'Actual')
trace1 = go.Scatter(x = X_train.T[0],y = lm.predict(X_train).T,mode = 'lines',name = 'Predicted')
adanient_data = [trace0,trace1]
layout.xaxis.title.text = 'Day'
plot2 = go.Figure(data=adanient_data, layout=layout)
```

```
[69]: iplot(plot2)
```

```
[70]: #Calculate scores for model evaluation
scores = f'''
{'Metric'.ljust(10)}{'Train'.center(20)}{'Test'.center(20)}
{'r2_score'.ljust(10)}{r2_score(Y_train, lm.
    predict(X_train))}\t{r2_score(Y_test, lm.predict(X_test))}
{'MSE'.ljust(10)}{mse(Y_train, lm.predict(X_train))}\t{mse(Y_test, lm.
    predict(X_test))}
'''
print(scores)
```

Metric	Train	Test
r2_score	0.8217821409507651	0.8362676928379684
MSE	160795.17281890797	158270.4011447389

## 1.8 Conclusion:

- We can observe that the accuracy achieved by the state-of-the-art ML model is no better than simply guessing with a probability of 50%. Possible reasons for this may be the lack of data or using a very simple model to perform such a complex task as Stock Market prediction.

```
[ ]:
```